

Infra

Imahn Shekhzadeh
`imahn.shekhzadeh@posteo.de`

March 2, 2025

Contents

1. Bash & Linux	7
1.1. Pattern Matching	7
1.2. Multi-Line Comments	7
1.3. File Download	7
1.4. for-loops	7
1.5. Argument Retrieval	7
1.6. Colored Outputs	8
1.7. String/File/Directory Operations	8
1.8. Monitoring	9
1.9. Systems Information	9
1.10. Port Number	10
1.11. CUDA	10
1.12. tmux	10
2. Docker	11
2.1. Installation	11
2.2. Basics	11
2.3. Dockerfile	11
2.4. Docker images	13
2.5. Docker containers	13
2.5.1. Basics	13
2.5.2. Passing Arguments	15
2.5.3. Listing & Stopping	16
2.6. Storage	16
3. AWS S3	17
3.1. Installation & Configuration	17
3.2. AWS Credentials (Profiles)	18
3.3. Buckets	18
3.3.1. Creation	18
3.3.2. Listings	18
3.3.3. File Copying	19
3.3.4. Directory Copying	19
3.3.5. Directory/File Deletion	19
3.4. Cloudpathlib	19
4. Conda	21
4.1. Installation of Environments	21
4.2. Export	21
4.3. Installation & Removal of Packages	22
4.4. Usage in VSCode	22

5. Git	23
5.1. <code>init</code>	23
5.2. <code>config</code>	23
5.3. Branches	23
5.3.1. Creation	23
5.3.2. Deletion	23
5.3.3. Listing	24
5.3.4. Merging	24
5.4. Stashing	26
5.5. Checking History	26
5.6. <code>rm</code>	26
5.7. <code>filter-repo</code>	27
5.7.1. Installation	27
5.7.2. Usage	27
5.8. <code>restore</code>	27
5.9. <code>reset</code>	28
5.10. <code>clean</code>	28
5.11. <code>blame</code>	28
5.12. <code>log</code>	28
5.13. Repo Renaming	29
5.14. Remote URL	29
5.15. Repo Change	29
5.16. Remote Repo Creation	29
5.17. Pull Requests	29
5.18. Updating Files from other Branch	30
6. Remote Development	31
6.1. Connection	31
6.2. Troubleshooting	31
7. Java	33
7.1. <code>ArrayList</code>	33
7.2. Classes & Interfaces	33
7.2.1. Basics	33
7.2.2. <code>instanceof</code>	35
7.2.3. <code>sealed</code>	36
7.3. Commenting	37
7.4. <code>JavaDoc</code>	37
7.4.1. <code>HTML</code>	37
7.4.2. <code>Tags</code>	37
8. Python	41
8.1. Config File & JSON Files	41
8.2. Jupyter Notebooks	41
8.3. <code>Map</code> (Built-In Function)	42
8.4. <code>CPU Brand</code>	42
8.5. <code>Inheritance</code>	43
9. PyTorch	45
9.1. <code>Leaf Tensors</code>	45

9.2. Autograd & Backward	46
9.3. Half-Precision	46
9.4. Miscellaneous	47
10. Jax	49
A. .bashrc	51
B. Amazing Programs, Extensions, Plugins & Packages	57
C. Opening Programs from the CLI in Linux	59
D. VSCode	61
D.1. Recommended Extensions	61
D.2. Debugging	61
D.3. <code>settings.json</code>	61
D.4. Fix Unresolved Python Imports	61
D.5. Opening a Duplicate Workspace	61
D.6. <code>settings.json</code>	62

Note

The commands in this document might only run through if you use the `.bashrc` file provided in App. A

1. Bash & Linux

1.1. Pattern Matching

In Bash, `[[]]` over `[]` is preferred, since `[[]]` is safer and more capable within Bash scripts. Within `[]`, where word splitting and filename expansion do occur, it is good practice to double-quote variables. But it is safe to omit the double-quotes for e.g. `##` within `[[]]`.

1.2. Multi-Line Comments

Use : ' [...] ' for a multi-line comment.

```
1  : '  
2  This is a  
3  very long  
4  comment  
5  '
```

1.3. File Download

Downloading file from URL and allowing for redirects,

```
1  curl -Lo output.out https://url.com
```

When taking a GitHub link, note that you need to take the URL of the raw file.

1.4. for-loops

For this directory structure,

```
1  infra_upd.tex  
2  infra_upd.pdf
```

rename via

```
1  for file in infra_upd.*; do mv "$file" "${file/infra_upd/infra}"; done
```

What happens is a [substring replacement](#).

1.5. Argument Retrieval

Retrieving all but the first argument,

```
1 bash_func(){
2     shift
3
4     echo "all provided args (except the first): $@"
5 }
```

Doing this $N \geq 1$ -times,

```
1 test_sth(){
2     shift
3     ...
4     shift
5
6     echo "all provided args (except the first N): $@"
7 }
```

1.6. Colored Outputs

Using colored outputs in Bash, cf. `str_diff` in App. A.

Personally, I use the following color scheme for the CLI,

1. monokai color scheme, i.e. dark gray background (#272822) with light peach color for the text (#F8F8F2).
2. File paths are still displayed in blue, which is suboptimal, to change the color to the better readable cyan-blue color, click on the three horizontal lines in the CLI, then on **Preferences**, then choose the currently active color, switch to the **Colors** tab, then go to **Palette**, click on the blue color & instead use the color #66D9EF

where `-e` stands for human readability and `-s` for summarizing.

1.7. String/File/Directory Operations

Appending line to file (`-a`: appending, otherwise `tee` overwrites `output.out` if existent),

```
1 echo "this is a line" | tee -a output.out # -a: appending, important
```

Checking whether string is empty,

```
1 [[ -z "$env_name" ]] && echo "The string is empty."
```

Finding out size of file/directory,

```
1 du -hs <path_to_file_or_dir> # du -hs file.ext
2
3 # for shorter summary (single quotation strings required)
4 du -hs <path_to_file_or_dir> | awk '{print $1}'
```

Unzipping a file via the CLI,

```
1 unzip /path/to/file.zip -d /path/to/destination
```


Zippping a directory via the CLI,

```
1 zip -r archive_name.zip /path/to/directory
```

and zipping several files,

```
1 zip archive_name.zip /path/to/file1 /path/to/file2 /path/to/file3
```

Opening a file and automatically scrolling to the bottom,

```
1 less +G /path/to/file.ext
```

Searching for files with specific extension, e.g. `.ext`:

```
1 find . -name "*.ext"
2 # find . -name "*.png"
```

Creating new directory including all parent directories (`-p` option is safe, since if directory is already existent, no error will be outputted),

```
1 mkdir -p <dir>
```

Comparing the contents of two directories,

```
1 diff -r --color directory1 directory2 # `-r` for recursive comparison
2 diff -rq --color directory1 directory2 # `-q` suppresses the output of
   differences and only shows which files differ
```

Ignoring files only existent in one of the directories (which treats absent files as empty),

```
1 diff -rq --color --unidirectional-new-file directory1 directory2
```

1.8. Monitoring

```
1 htop
```

RAM usage,

```
1 free -h
```

1.9. Systems Information

Retrieving the number of available CPU resources,

```
1 echo "$(nproc)"
```

Print day and time from CLI,

```
1 echo "$(date +%d_%m_%y-%H_%M_%S)"
2 # echo "$(date +%dp%mp%y-%Hp%Mp%S)"
```

Listing all available kernels in Debian-based Linux systems,

```
1 dpkg --get-architecture | grep linux-image
```

Currently active kernel version,

```
1 uname -a
```

1.10. Port Number

To check whether port is used or not,

```
1 ss -tuln | grep :<port_number>
2 # e.g.: `ss -tuln | grep :11434`
```

1.11. CUDA

- When you need to find out the CUDA version installed, install `nvidia-cuda-toolkit`, but do NOT reboot. After its use, immediately remove this package and any package installed alongside with it!
- In case NVIDIA drivers do not allow for boot into Ubuntu, e.g. because you did not uninstall the `nvidia-cuda-toolkit` package,
 1. Boot into an older kernel version of Linux (in order to get there, do a "hard" reboot, and then go into "Advanced options for Ubuntu", and choose an older kernel version).
 2. Once booted into the older kernel version, I removed 'nvidia-cuda-toolkit' and rebooted.
 3. After a few more hard reboots and booting into the older kernel version, at some point, the newer kernel version was picked up and worked again.
 4. Now to fix the monitors (because dual-monitor setup didn't work), I had to open the program "Additional Drivers" and change the driver from the open-source version to an NVIDIA proprietary one.
 5. Then I had to install CUDA according to these instructions.
 6. For PyTorch to recognize the GPU, I had to reboot.

1.12. tmux

- Creating a new session:

```
1 tmux new -s <name>
2 # tmux new -s mysession
```

- To detach from the current TMUX session and leave it running in the background, press CTRL + B and then type D.
- Reconnecting to a TMUX session,

```
1 tmux attach -t <name>
2 # tmux attach -t mysession
```

- Killing a TMUX session from outside,

```
1 tmux kill-session -t <name>
2 # tmux kill-session -t mysession
```

2. Docker

2.1. Installation

- Follow this great tutorial by DigitalOcean.
- To start the Docker daemon,

```
1 sudo systemctl start docker
```

- To use NVIDIA GPUs (both in PyTorch & Jax), install the NVIDIA Container Toolkit
- Once done with the installation of the NVIDIA Container Toolkit, proceed with the configuration. During the configuration, it will be necessary to restart the docker daemon, which you can achieve as follows:

```
1 sudo systemctl restart docker
```

- Checking that the docker installation was successful,

```
1 docker run hello-world
```

2.2. Basics

- Interactive start of containers:

```
1 d ps -a # find out ID (also docker container name)
2 d start -i ID
```

- Copying files from local system to docker container and vice versa; **run both commands from local CLI**

```
1 d cp file_name container_ID:/target_dir # local -> docker
2 d cp container_ID:/file_name dir_name # docker -> local
```

2.3. Dockerfile

- When you find the command for pulling a docker image on <https://hub.docker.com>, e.g.

```
1 d pull ubuntu:jammy-20231004
```

then in the Dockerfile, just write

```
1 FROM ubuntu:jammy-20231004
```

2. Docker

When no tag is specified, by default the *latest* one will be taken. However, using the *latest* tag can potentially cause issues with reproducibility and consistency, because you might pull a different version of the image at different times without knowing it if the latest tag gets updated. **For more predictable builds, it is advised to use a specific version tag.**

- Note that the structure of the *docker pull* command is

```
1 d pull [OPTIONS] NAME[:TAG|@DIGEST]
```

In general, the *NAME* is in the format *repository/image*. If *repository* is not specified, Docker assumes the image is located in the default DockerHub library repository. However, many images (like PyTorch) are hosted under a specific user or organization's namespace on DockerHub, rather than the top-level library. That's why the command for the docker pull (for the latest tag) reads

```
1 d pull pytorch/pytorch
```

- If using a Docker image like *pytorch/pytorch:latest*, conda is already installed. In this case, the default environment is named *base*, which is a common practice in Docker images with conda – unless otherwise stated.
- Copying local scripts into docker container,

```
1 COPY relative/path/to/script.py .
```

From the documentation:

Multiple `<src>` resources may be specified but the paths of files and directories will be interpreted as relative to the source of the context of the build.

It is also important to put the `.` at the end, since it represents the destination in the Docker image where the file should be copied. The dot `.` refers to the current working directory inside the Docker image, which is determined by the `WORKDIR` command in the Dockerfile. If `WORKDIR` is not set, it defaults to the root directory (`/`) of the image.

Also, each time the script `relative/path/to/script.py` changes, the Dockerfile needs to be rebuilt — **however, a cached version will be used, which speeds things up.**

- Copying local dirs into docker container,

```
1 COPY relative/path/to/dir/ .
```

- Running a Dockerfile,

```
1 d build -f file_name -t img_name .
2 d build -f file_name -t img_name:tag_name . # tag name optional, but
      recommended, e.g. 1.0 (no quotes required)
```

where `img_name` will be the name of the newly created image, `tag_name` the tag name and `file_name` the name of the docker file.

The `.` specifies the context of the build, which is the current directory in this case. If `custom_docker_file` is in another path, it can be easily provided,

```
1 d build -f file_name -t img_name:tag_name /path/to/build/context
```

- In order to forcefully rebuild the Dockerfile from scratch, i.e. without the use of a cached version, use the `--no-cache` option,

```
1 d build -f file_name --no-cache -t img_name:tag_name .
```

- Via

```
1 EXPOSE custom-port-number
2 # EXPOSE 80
```

it is possible to expose a port. Note that port exposure is related to network access. Note that even though network access might not be needed, there is still no harm in exposing a port (since an exposure of the port does not make the docker container more vulnerable).

2.4. Docker images

- Check all available Docker images via

```
1 d images
```

- Cleaning up dangling docker images — these are the entries with `<none>` in the repository or tag name in the output of the previous algo:

```
1 d image prune -f
```

If the total acclaimed disk space is 0B, give this command a try, which can help clean up *a lot of* disk space:

```
1 d system prune
```

- Removing a Docker image — **only do this when finished with using the image**

```
1 d image rm img_name:tag_name
2 # d container rm <container_id> # in case some containers are using
   the image
```

2.5. Docker containers

2.5.1. Basics

- In case of shared memory issues in the docker container, the shared memory — which is by default 64 MB — needs to be increased when the docker container is run,

```
1 d run --shm-size 512m [...] # requesting memory in MB
2 # d run --shm-size 1G [...] # requesting memory in GB
```

2. Docker

- Running Docker images – without being able to utilize NVIDIA GPUs:

```
1 d run -it img_name # if `tag_name` was not provided
2 d run -it img_name:tag_name # if `tag_name` was provided during
    build (recommended)
```

- Running Docker images & utilizing GPUs:

```
1 d run --gpus all -it img_name
2 d run --gpus all -it img_name:tag_name # recommended
```

- To mount a local file to the container at runtime, do

```
1 d run -v /absolute/path/to/script.py:/path/to/workdir/script.py --
    gpus all -it img_name
2 d run -v /absolute/path/to/script.py:/path/to/workdir/script.py --
    gpus all -it img_name:tag_name # recommended, provide `img_name`
    & `tag_name`
```

The mounting expects **absolute** file paths on the side of the host machine.

- Note that you can include the bash command **pwd** to avoid having to manually pass absolute paths for the mounting

```
1 d run -v $(pwd)/script.py:/path/to/workdir/script.py --gpus all -it
    img_name:tag_name # recommended, provide `img_name` & `tag_name`
```

If you need the container to reflect changes made to the scripts on the host without rebuilding the image every time, you would use the `-v` flag to mount the directory. If the scripts won't change, or you don't need to reflect changes in real-time, you don't need to mount the directory, as the necessary scripts have already been copied into the image during the build process.

- It is also possible to directly mount directories:

```
1 d run -v $(pwd)/dir_path:/path/to/workdir --gpus all -it img_name:
    tag_name
```

Note that the specified directory from the host is mounted into the container at the specified mount point. If there are any existing files or directories in the container at the mount point, they become obscured by the mount.

- In several cases it can be useful to remove the docker container right after execution: When you...
 - ...are running many short-lived containers, like during development or testing,
 - ...want to avoid manual cleanup of stopped containers later on,
 - ...are running containers for one-off tasks that do not need to persist any state after they are finished.

In this case,

```
1 d run --rm -v $(pwd)/dir_path:/path/to/workdir --gpus all -it
    img_name:tag_name
```

- It is also possible to mount two separate host directories to two separate directories within the container,

```
1 d run --rm -v $(pwd)/dir_path1:/path/to/workdir1 -v $(pwd)/dir_path2
   :/path/to/workdir2 --gpus all -it img_name:tag_name
```

This will not cause any overwriting as each `-v` flag creates a unique mount point inside the container.

- Finding out the python version of the Docker image

```
1 d run -it --rm img_name:tag_name python3 --version
```

This command will immediately remove the container after execution.

- It is also possible to interact with a docker container,

```
1 docker run -it --rm img_name:tag_name /bin/bash
```

2.5.2. Passing Arguments

It is possible to pass arguments when running a docker container.

1. Assuming you have a bash script `run_scripts.sh`, in which a Python script, e.g.

```
1 #!/bin/sh
2 isort /app/scripts/*.py
3 black /app/scripts/*.py
4
5 python3 -B /app/scripts/test_script.py
6 python3 -B /app/scripts/test_anil.py
```

Modify this bash script s.t. any arguments passed to the CLI when running the docker container are picked up,

```
1 python3 -B /app/scripts/test_anil.py "$@"
2 # python3 -B /app/scripts/test_script.py "$@" # alternative
```

2. Rebuild (!) the docker image.
3. Now run the docker container as follows:

```
1 d run --rm -v $(pwd)/dir_path:/path/to/workdir --gpus all -it
   img_name:tag_name arg1 arg2
2 # d run --rm -v $(pwd)/dir_path:/path/to/workdir --gpus all -it
   img_name:tag_name --n_ways 1 --k_shots 1 # example
```

2.5.3. Listing & Stopping

- Listing all running containers,

```
1 d ps
```

Listing only the container ID of all running containers,

```
1 d ps -q
```

- Stopping a running container,

```
1 d stop container-ID
```

- Stopping a running container and removing it,

```
1 d stop container-ID && d rm container-ID
```

2.6. Storage

- In case you want `docker` to install images and containers in a separate drive, e.g. one under `/media/user-name/samsung_500`, you can follow these steps:

1. Stop the docker service,

```
1 sudo systemctl stop docker
```

2. Move the directory `/var/lib/docker` to the separate drive,

```
1 sudo mv /var/lib/docker /media/user-name/samsung_500/docker
```

Note that you should *not* create the `docker` directory directly under `/media/user-name/samsu`

3. Configure `docker` to use the new directory by editing the `docker` daemon configuration file,

```
1 sudo nano /etc/docker/daemon.json
```

and then adding the following configuration,

```
1 {
2   "data-root": "/media/user-name/docker"
3 }
```

4. Restart the `docker` service,

```
1 sudo systemctl restart docker
```

5. Follow the instructions from the NVIDIA Container Toolkit configuration to ensure that NVIDIA GPUs can still be used inside the containers.

3. AWS S3

3.1. Installation & Configuration

1. Installation instructions
2. The CLI will display the path under which the *aws* package was installed, but it might be sufficient to simply run

```
1  aws
```

Double check by running

```
1  which aws
```

3. After installation, configuration is necessary. For this run

```
1  aws configure
```

You can leave these fields empty:

```
1  Default region name [None]:  
2  Default output format [None]:
```

A configuration file will be saved under

```
1  ~/.aws/credentials
```

4. In the case you are a member of UNIGE, you can obtain the AWS access key ID and the secret access key as follows:

```
1  echo -n "$user_name" | base64 # the `-n` is important in this  
    context  
2  echo -n "$passwd" | md5sum
```

where `$user_name` and `$passwd` need to be provided

Otherwise, you need login to the AWS Management Console.

5. To test the configuration was successful, do this:

```
1  aws s3 ls --endpoint-url https://your-custom-s3-endpoint.com
```

where you replace the endpoint-url `https://your-custom-s3-endpoint.com` with yours.

3. AWS S3

3.2. AWS Credentials (Profiles)

- It is possible to use several profiles in the file `~/.aws/credentials`.
- For example,

```
1 [default]
2 aws_access_key_id = YOUR_DEFAULT_ACCESS_KEY
3 aws_secret_access_key = YOUR_DEFAULT_SECRET_KEY
4
5 [profile1]
6 aws_access_key_id = ANOTHER_ACCESS_KEY_ID
7 aws_secret_access_key = ANOTHER_SECRET_ACCESS_KEY
8
9 [profile2]
10 aws_access_key_id = YET_ANOTHER_ACCESS_KEY_ID
11 aws_secret_access_key = YET_ANOTHER_SECRET_ACCESS_KEY
```

Using specific profile when running `aws cli` commands via `--profile` option in the command:

```
1 aws s3 --profile profile1 [...]
2 # aws s3 --profile default [...]
```

3.3. Buckets

One can have several buckets.

3.3.1. Creation

- Creating a new bucket,

```
1 aws s3api create-bucket --bucket custom-bucket-name --endpoint-url
  https://custom-s3-endpoint.com --profile default
```

3.3.2. Listings

- Directly “folder” contents of an s3 bucket,

```
1 aws s3 ls s3://custom-bucket-name --recursive --endpoint-url https
  ://custom-s3-endpoint.com --profile default # `--recursive`
  optional
```

- Showing file contents,

```
1 aws s3 ls s3://custom-bucket-name/prefix/ --recursive --endpoint-url
  https://custom-s3-endpoint.com --profile default # `--recursive`
  optional
```

Note that the `/` at the end of the prefix (“folder”) is necessary.

3.3.3. File Copying

- Local machine → S3:

```
1  aws s3 cp path/to/custom_file.ext s3://custom-bucket-name/path/to/
   custom_file.ext --endpoint-url https://custom-s3-endpoint.com --
   profile default
```

- S3 → local machine:

```
1  aws s3 cp s3://custom-bucket-name/path/to/s3_file.ext custom/
   destination --endpoint-url https://custom-s3-endpoint.com --
   profile default
```

3.3.4. Directory Copying

- Local machine → S3:

```
1  aws s3 sync path/to/dir s3://custom-bucket-name/path/to --endpoint-
   url
2  https://custom-s3-endpoint.com --profile default
```

3.3.5. Directory/File Deletion

- Deleting a folder (which is essentially a prefix in S3) and its contents in an S3 bucket,

```
1  aws s3 rm s3://your-bucket-name/path-to-your-folder --recursive --
   endpoint-url https://custom-s3-endpoint.com --profile default
```

- Deleting a file,

```
1  aws s3 rm s3://your-bucket-name/path-to-your-file.out --recursive --
   endpoint-url https://custom-s3-endpoint.com --profile default
```

3.4. Cloudpathlib

- When you use the cloudpathlib module, and you want to specify a profile, do this:

```
1  from cloudpathlib import S3Path, S3Client
2
3  # Create an S3 client with a specific AWS profile
4  s3_client = S3Client(
5      aws_access_key_id=aws_access_key_id,
6      aws_secret_access_key=aws_secret_access_key,
7      endpoint_url=endpoint_url,
8      profile_name="profile1", # specify profile here
9  )
10
11 # Make `client` default:
12 client.set_as_default_client()
```


4. Conda

- Retrieving information about currently activated conda environment,

```
1 conda info
```

- Listing all installed environments,

```
1 conda env list
```

4.1. Installation of Environments

- Installing conda with specific python version,

```
1 # only `myenv` needs to be specified (quotation marks necessary)
2 env_name="myenv" && conda create -n "$env_name" python=3.11.3 -y &&
  conda activate "$env_name"
```

As of Oct 16, I wouldn't recommend installing python 3.12.0 yet — I got a lot of unmet dependency problems when trying to install torch 2.1 with NVIDIA Cuda version 11.8 afterwards.

- Installation of conda environment from bash file:

```
1 conda deactivate # go into base environment
2 source conda/filename.sh
3 touch .env
```

- Completely remove conda environment,

```
1 conda deactivate && conda remove -n custom-env-name --all -y
```

4.2. Export

- Exporting an .yaml-file to share with others for reproducibility,

```
1 conda env export > environment.yaml
```

- Line “Prefix:” at end of .yaml file can be safely deleted, for details cf. [here](#)

4.3. Installation & Removal of Packages

- Installation of packages from `pyproject.toml` file,

```
1 pip install -e .
```

If there is not enough free space, do

```
1 TMPDIR=[...] pip install -e .
```

where `TMPDIR` needs to exist.

If you want to install a specific version of a package via `pip` — and you do not use a `pyproject.toml` file — do

```
1 pip install package==version
```

- Installing specific conda package in a specific version — note that the specification of a version number is optional,

```
1 conda install -c conda-forge custom-pkg-name=version-number -y
2 # conda install -c conda-forge cloudpathlib=0.15.1 -y
3 # conda install -c conda-forge cloudpathlib -y
```

- Removing list of packages from conda environment,

```
1 conda remove -n custom-env-name pkg1 pkg2 ... pkgN -y
2 # conda remove -n google_jax matplotlib -y
```

4.4. Usage in VSCode

- Selecting a conda environment in VSCode, do `Ctrl + Shift + P` and type `Python: select interpreter`.

5. Git

5.1. init

Initialize a new repository in the current working directory; creates a new, empty directory:

```
1 git init
```

To check whether the repository has been initialized, look out for the hidden file `.git`.

Note that this repository will not yet be visible remotely, e.g. on GitLab or GitHub!

5.2. config

Setting your username and e-mail:

```
1 git config --global user.name "Blub blub"
2 git config --global user.email "blub@blub.com"
```

The `--global` flag sets the configuration for all repositories on your machine. If you want to set the username and email for a specific repository only, omit the `--global` flag and run the commands inside that repository.

To check that these settings were successful, run

```
1 git config --global --list
2 git config --list
```

5.3. Branches

5.3.1. Creation

```
1 git switch -c <new_bname>
2 git push -u origin <new_bname>
```

5.3.2. Deletion

Delete a remote branch via

```
1 git push origin --delete <branch_name>
```

To delete a local branch, first switch to another branch and then do

```
1 git branch -d <branch_name> # if branch has already been merged into
    upstream branch
2 git branch -D <branch_name> # force deletion; irrespective of merge
    status (equivalent to `git branch --delete --force <branch_name>`)
```

5. Git

If a branch is deleted remotely, it might still exist locally.

When merging a branch into another one remotely, e.g. via GitLab, and deleting it in the process, it can happen that `git branch -a` would indicate that the branch still exists. To correct this,

```
1 git fetch --prune
```

The branch might still exist locally, just run Alg. 5.3.2.


5.3.3. Listing

List all remote and local branches:

```
1 git branch -a
```


5.3.4. Merging

Basics




When merging `main` and `new_branch`, the commits D and D' ' will either be auto-merged or there will be a merge conflict. In case the merge is successful or the merge conflicts have been resolved, there will be a new **merge commit** in the `main` branch. In case there was an auto-merge, git automatically creates the merge commit, otherwise the user needs to do this after resolving the merge conflicts.

For this commit history,



merging `main` with `new_feature` would result in a fast-forward (and no merge commit is created),



Merging `new_feature` into `main` can be done via

```
1 git switch main && git merge new_feature
```

Removing the commit message Merge pull request #6 from [...],

```
1 git switch main && git merge new_feature --log
```

Running dry merges to proactively check for conflicts *if* a merge was performed,


```
1 git merge --no-commit --no-ff branch-name && git merge --abort
```

For a fast-forward only,

```
1 git merge branch_to_be_merged --ff-only
```

To abort an ongoing merge,

```
1 git merge --abort
```

When merging a branch into another one remotely, e.g. via GitLab, and deleting it in the process, it can happen that `git branch -a` would indicate that the branch still exists. To correct this, cf. Alg. 5.3.2.

Conflicts

1. Resolving a merge conflict,

```
1 git mergetool
```

2. Confirm with Enter that you want to use vimdiff as default editing tool. vimdiff display will resemble the following structure:

```
1 | LOCAL | BASE | REMOTE |
2 | MERGED
```

If file did not already exist in BASE, then we need this view:

```
1 | LOCAL | MERGED | REMOTE |
```

LOCAL – Current branch

BASE – Common ancestor (how did the file look like before both changes?)

REMOTE – File that I am merging into the current branch

MERGED – Merge result

3. It is probably easiest to take the merged view and edit it directly. In the vim editor, an entire line can be deleted by pressing D (no control before!). If I instead wanted the changes from either LOCAL, BASE or REMOTE, you have to do one of these,

```
1 :diffg LO
2 :diffg BA
3 :diffg LO
```

Of course, the merged view can also be edited directly.

4. Type

```
1 :wqa
```

into vim. Afterwards, do not forget to commit and push. And if you want, do

```
1 git clean -f
```

5.4. Stashing

The stash follows a FIFO (first-in, first-out) principle.

To include untracked files for the stash, do

```
1 git stash --include-untracked
```

In case you want to stash only one file, do (<https://stackoverflow.com/a/3041055/15528750>):

```
1 git stash push <file_name> --include-untracked
```

When applying the stash (to make sure that the files are staged also after applying the stash, use `-index` option):

```
1 git stash apply stash@{0} --index
```

To remove a particular item from the stash, do (first item starts with 0)

```
1 git stash drop stash@{2}
```

The last two commands can also be combined. To find out the index number on the stash, we can simply do

```
1 git stash list
```

To remove all stashes, do

```
1 git stash clear
```

This command is irreversible, though! To remove a specific stash in git, you can use the `git stash drop` command followed by the stash identifier, which is typically in the form `stash@<stash_number>`. For example, if you want to remove the stash with the identifier `stash@2`, you would run:

```
1 git stash drop stash@{2}
```

5.5. Checking History

- Viewing the history of commits,

```
1 git log
```

- Viewing a specific file,

```
1 git show <commit-hash>:<file-name>  
2 # git show 123abc:example.txt
```

5.6. rm

- To remove a file/folder that is already tracked, adding it to `.gitignore` won't remove it (though this also needs to happen). For this, do:

```
1 git rm --cached <file>  
2 # git rm -r --cached <folder>  
3 git push origin <bname>
```

- Adding the file/folder to *.gitignore* is still a good idea, though, since the file/dir won't be removed locally with the commands.
- Note that if you also want to delete the file from the history, then you should use `git filter-repo`, cf. Sec. 5.7.

5.7. filter-repo

5.7.1. Installation

To delete a file from the history, first install `filter-repo`:

```
1 sudo apt-get install git-filter-repo -y
```

5.7.2. Usage

To rewrite the history:

```
1 git filter-repo --path PATH_TO_FILE --invert-paths --force
```

If `PATH_TO_FILE` still exists remotely, then add the remote repository:

```
1 git remote add origin git@github.com:username/repository.git
```

And then verify this was successful via

```
1 git remote -v
```

Now force-push.

5.8. restore

Restoring specific file to state of any previous commit,

```
1 git restore --source=<commit-hash> <file-path>
2 # git restore --source=HEAD README.md
3 # git restore --source=HEAD .
```

To restore the content in the index, additionally use `--staged`:

```
1 git restore --staged --source=<commit-hash> <file-path>
2 # git restore --staged --source=HEAD README.md
```

Only restore the index if you know what the implications are.

To restore both the index and the working tree, add `--staged --worktree`:

```
1 git restore --staged --worktree --source=<commit-hash> <file-path>
2 # git restore --staged --worktree --source=HEAD README.md
```

Only restore the worktree if you know what the implications are.

If you want any of the changes to be permanent, type `git push` afterwards.

5.9. reset

If you made a commit that is not pushed yet, and you do not want to throw away the commit, but just continue working on it, then do

```
1 git reset HEAD~1
```

If you want to throw the last n commits, then do

```
1 git reset --hard HEAD~n
```

If the commits were already pushed, then force push.

5.10. clean

Note that neither Alg. 5.9 nor Alg. 5.9 will remove untracked files or directories. For this, do

```
1 git clean -fd
```

For this, please note that the untracked files will be irrevocably deleted!

5.11. blame

The command `git blame` can be used to understand which person authored a revision for each line of a file. Basic usage:

```
1 git blame <filename>
```

Blaming a **specific commit**:

```
1 git blame <commit_hash> <filename>
```

Ignoring whitespace changes:

```
1 git blame -w <filename>
2 git blame -w <commit_hash> <filename>
```

5.12. log

There are different ways to format the output. In order to get a graph-like visualization of the branches and commits, one can use:

```
1 git log --oneline --all --graph --decorate
```

This command is also very useful to see the origin of branches. If you want more detailed information, e.g. timestamps and author information, remove the `--oneline`:

```
1 git log --all --graph --decorate
```

5.13. Repo Renaming

```
1.
1  gh repo rename <new-repo-name>
2  # gh repo rename new-repo-name
```

```
2.
1  git remote set-url origin <new-repo-url>
2  # git remote set-url origin https://github.com/username/new-repo-
   name.git
```

double-check via

```
1  git remote -v
```

which lists the remote names and their URLs. No force push or the alike is necessary for the changes to take place.

5.14. Remote URL

Obtaining the remote URL,

```
1  git remote get-url origin
2  git remote get-url origin | sed 's/\.git$//' # optional: trim output
```

5.15. Repo Change

Moving all files from `branch-to-move` to `branch-to-merge-into` and preserving the commit history — do all of this while in the old repo,

```
1  git remote -v # check existing remotes
2  git remote add <target> https://target-repo-url.git # add new remote
3  # git remote add new-remote url
4  git push target branch-to-move:branch-to-merge-into
```

Do all of this in the old repo. If issue emerges during the last step, reclone the new repo and check whether this solves the issue.

5.16. Remote Repo Creation

Then install GitHub CLI and do

```
1  gh repo create <repository-name> --public # --private
```

Then commit and push.

5.17. Pull Requests

5. Git

```
1 gh pr create --base main --head "$bname" --title "Pin isort & black
2 versions" --body "This pull request fixes the issue that workflows
   fail
3 because of different isort/black versions used in the workflows &
4 specified in the \`pyproject.toml\` file."
```

Note that `bname` is a bash function defined in App. A. Escaping the `\` is necessary, since in shell commands, backticks (```) are used to execute commands and substitute their output into the command line.

5.18. Updating Files from other Branch

When working on `branch_my`, it is possible to incorporate changes from another branch `branch_x`,

```
1 git switch <branch_x>
2 git pull origin <branch_x>
3 git switch <branch_my>
4 git rebase -i <branch_x>
5 git push origin <branch_my> --force
```

Save the interactive view via `:wq` and make a force push to `branch_my`.

6. Remote Development

6.1. Connection

1. When connecting two machines remotely, install this extension on local machine (also directly in VSCode possible),
2. open VSCode on local machine,
3. press F1-button, choose “Remote-SSH: Connect to Host...” and type for the SSH host (optionally save it in the SSH config file) the same as in Algo. (B),
4. enter the passwd for the remote SSH host.

6.2. Troubleshooting

If you find you are getting a permission error for saving a file on the remote machine (in VSCode when doing the local coding), try

```
1 sudo chown custom-username path/to/custom/script.ext
```

`custom-username` here refers to the username on the remote machine. If the remote connection hung up,

```
1 fusermount -zu /path/to/dir
```


7. Java

7.1. ArrayList

- In Java, ArrayList objects can be initialized directly:

```
1 import java.util.ArrayList;
2
3 ArrayList<Integer> numbers = new ArrayList<Integer>(  
4     List.of(1, 2, 4, -5, 10)  
5 );
```

This only works in Java version 9+.

- From Java version 7+, you can use the **diamond operator** to omit the generic type on the RHS of the declaration:

```
1 import java.util.ArrayList;
2
3 ArrayList<Integer> numbers = new ArrayList<>(  
4     List.of(1, 2, 4, -5, 10)  
5 );
```

- To compare two ArrayList objects for equality, do **not** use the == operator; instead

```
1 import java.util.ArrayList;
2
3 ArrayList<Integer> numbers1 = new ArrayList<>(  
4     List.of(1, 2, 4, -5, 10)  
5 );  
6 ArrayList<Integer> numbers2 = new ArrayList<>(  
7     List.of(1, 2, 4, -5, 10)  
8 );  
9  
10 boolean equalValues = numbers1.equals(numbers2); // `true`  
11 boolean equalReferences = (numbers1 == numbers2); // `false`
```

7.2. Classes & Interfaces

7.2.1. Basics

- Classes, abstract classes and interfaces can be either public or package-private.
- Abstract classes can have constructors; interfaces, on the other hand, cannot.

7. Java

- When inheriting from a class, the subclass must implement the constructor if the super class has a non-default constructor and call the constructor from the super class.
- Abstract classes can contain public, protected, package private and private fields. In interfaces, fields are always **public static final**, i.e. there are no instance fields.
- All non-static methods in interfaces are implicitly abstract and public¹; however, static methods in interfaces can be concrete. Also, static methods in interfaces can be either public (which they are implicitly, cf. here²) or private, but not protected or package private (since package private would be specified by not adding any access level modifier, but in interfaces, methods are implicitly public).³
- In abstract classes, abstract methods need to be declared as such (via the **abstract** keyword) and they can be public, protected or package private (but not private!). Abstract classes can also have non-abstract (concrete) methods, which can be public, protected, package private or private⁴.
- An abstract class can implement an interface, yet does not have to implement the abstract methods; this has to be done by subclasses inheriting from the abstract class.
- When implementing an interface's method or an abstract method from an abstract class, it is recommended to use the **@Override** tag for better readability and maintenance.
- When implementing a static method in a class, it can only refer to static fields.
- The **final** keyword can be used independently from **sealed** classes, cf. Sec. 7.2.3, and means that no class can inherit from a final class.
- In order to extend interfaces, it is possible to add **default** methods, which can be overridden in the classes implementing the interface.

```
1  public interface Truck
2  {
3      void accelerate();
4
5      void break();
6
7      default void fly(){};
8  }
9
10 public class MAN implements Truck
```

¹Hence the access modifiers do not need to be specified. When writing a class that implements an interface, the class cannot reduce the visibility of the methods from the interface, which means that those need to be public. Other additional methods can of course have any access level modifiers.

2

Every method declaration in the body of an interface is implicitly public (§6.6). It is permitted, but discouraged as a matter of style, to redundantly specify the public modifier for a method declaration in an interface.

³It is generally the case that static methods cannot be **protected**, since they belong to the type where they are defined.

⁴Obviously, private non-abstract methods cannot be overridden, and private abstract methods do not exist.

```

11 {
12     @Override
13     public void accelerate(){};
14
15     @Override
16     public void break(){};
17
18     @Override
19     public void fly()
20     {
21         //
22     }
23 }

```

- When inheriting from another class and overriding (not overloading) a method from the super class, you can use the `@Override` tag to make this clear. While this is not strictly necessary, it makes the code better readable and it is easier to maintain.

```

1  public class Account {
2      protected double savings;
3
4      public void withdraw(double amount) {
5          this.savings -= amount;
6      }
7  }
8
9  public class SavingAccount extends Account {
10
11      @Override
12      public void withdraw(double amount) {
13
14          if (amount > 0 && amount > this.savings) {
15              this.savings -= 1.01 * amount;
16          }
17          else {
18              super.withdraw(amount);
19          }
20      }
21  }

```

7.2.2. instanceof

When a class implements an interface, then an object of this class is also an instance of the interface.

```

1  public interface InvestmentCalculator {
2      [...]
3  }
4

```

7. Java

```
5 public class ExponentialInvestmentCalculator implements
    InvestmentCalculator{
6     [...]
7 }
8
9 // In a main function:
10 ExponentialInvestmentCalculator calculator = new
    ExponentialInvestmentCalculator();
11 boolean instanceCheck = calculator instanceof InvestmentCalculator; //
    `true`
```

With **polymorphism**, you can also make `calculator` an object, which has as type the interface `InvestmentCalculator`:

```
1 // In a main function:
2 InvestmentCalculator calculator = new ExponentialInvestmentCalculator()
    ;
3 boolean instanceCheck = calculator instanceof InvestmentCalculator; //
    `true`
```

7.2.3. sealed

It is possible to define **sealed** classes, **sealed** abstract classes and **sealed** interfaces in Java17+. The keyword has the effect that it restricts which other classes or interfaces may extend or implement them. For example:

```
1 abstract sealed class Daimler extends Car permits EKlasse, CKlasse
2 {
3     void bite(){};
4
5     void doNothing(){};
6 }
```

Here, the class `Daimler` inherits from the class `Car` and only allows the subclasses `EKlasse` and `CKlasse`, which will need to be defined.

The subclasses `EKlasse` and `CKlasse` can now either be **final**, i.e. no other class can inherit from them, **sealed**, or **non-sealed**, which means that any subclass can inherit from the class (note that a sealed class cannot prevent its permitted subclasses from doing this). For example:

```
1 final class EKlasse extends Daimler
2 {
3     @Override
4     void bite()
5     {
6         //
7     }
8 }
```

Also, according to the documentation, any permitted subclasses must be accessible by the sealed class at compile time.

7.3. Commenting

For multi-line comments, use `/* comment...*/`, but if you want to comment a method, use `/** doc string...*/`.

7.4. JavaDoc

7.4.1. HTML

It is possible to generate an HTML version of the docs. For this, do

```
1 javadoc -private -d doc `file1.java` `file2.java` [...]
2 # javadoc -private -d doc Account.java SavingsAccount.java
```

If you intend to publish the docs, then remove the `-private` option.

7.4.2. Tags

- `@see`: Reference other functions; these references are clickable. Example:

```
1 public class Account {
2     protected double savings;
3
4     public void withdraw(double amount) {
5         this.savings -= amount;
6     }
7 }
8
9 public class SavingAccount extends Account {
10
11     /**
12      * Saving account discourages money being withdrawn through a 1%
13      * relative penalty relative to the amount that is withdrawn.
14      *
15      * @param amount Amount to be withdrawn.
16      * @return
17      * @see Account#withdraw(double)
18      */
19     @Override
20     public void withdraw(double amount) {
21
22         if (amount > 0 && amount > this.savings) {
23             this.savings -= 1.01 * amount;
24         }
25     }
26 }
```

- `@throws`

```
1 public class Account {
2
```

7. Java

```
3      /**
4       * Withdraw money from account.
5       * @param amount Amount to be withdrawn.
6       * @return
7       * @throws IllegalArgumentException if the amount is negative
8       */
9      @Override
10     public void withdraw(double amount) {
11
12         if (amount < 0) {
13             throw new IllegalArgumentException(
14                 "Amount must be strictly positive!"
15             );
16         }
17     }
18
19 }
```

- `@deprecated`: Used to indicate that a method or class is deprecated.
- `@link`: Used to create hyperlinks in comments; references methods, constructors, fields or other classes in the documentation.

```
1     public class Account {
2         protected double savings;
3         protected double withdrawalPenalty;
4
5         public Account(double savings, double withdrawalPenalty) {
6             this.savings = savings;
7             this.withdrawalPenalty = withdrawalPenalty;
8         }
9
10        public void withdraw(double amount) {
11            this.savings -= amount;
12        }
13    }
14
15    public class SavingAccount extends Account {
16
17        public SavingAccount(double savings, double withdrawalPenalty) {
18            super(savings, withdrawalPenalty);
19        }
20
21        /**
22         * Saving account discourages money being withdrawn through a 1%
23         * relative {@link withdrawalPenalty penalty} relative to the
24         * amount that is withdrawn.
25
26         *
27         * @param amount Amount to be withdrawn.
28         * @return
29         */
30    }
```

```
26      * @see Account#withdraw(double)
27      */
28      @Override
29      public void withdraw(double amount) {
30
31          if (amount > 0 && amount > this.savings) {
32              this.savings -= ((1 + withdrawalPenalty) * amount);
33          }
34      }
35  }
```

Note that is not necessary to use the `@link` tag when using the `@see` tag, since `@see` automatically leads to a hyperlink.

8. Python

8.1. Config File & JSON Files

- When using `argparse` in combination with a JSON configuration file, the JSON keys need to match the long option names specified in `parser.add_argument()` method calls. The `argparse` module itself does not automatically recognize abbreviated forms from a JSON file.

8.2. Jupyter Notebooks

- Converting jupyter notebooks into PDFs,

```
1  for nb in /path/one/Notebook1.ipynb /path/two/Notebook2.ipynb [...]
2  do
3  jupyter nbconvert --to pdf "$nb"
4  done
```

Wildcarding notation would also work,

```
1  # optionally: `output_dir="[...]"`
2
3  for nb in *.ipynb; do
4  nb_name="${nb%.ipynb}"
5  jupyter nbconvert --to pdf "$nb" # `--output "$output_dir/$nb_name.
   pdf"`
6  done
```

- Changing default theme of notebooks,

```
1  conda install conda-forge::jupyterthemes
2  jt -l # get list of all available themes
3  jt -t <theme-name> # change theme
4  # jt -t onedork
```

Alternative installation via pip,

```
1  pip install jupyterthemes
```

- Displaying all statements in a Jupyter NB, e.g.,

```
1  from IPython.core.interactiveshell import InteractiveShell
2  InteractiveShell.ast_node_interactivity = "all"
3
4  x = 2
```

8. Python

```
5   y = 3
6
7   x
8   y
```

8.3. Map (Built-In Function)

- Function signature:

```
1   map(function, iterable, *iterables)
```

Description provided in the documentation:

Return an iterator that applies function to every item of iterable, yielding the results. If additional iterables arguments are passed, function must take that many arguments and is applied to the items from all iterables in parallel. With multiple iterables, the iterator stops when the shortest iterable is exhausted. For cases where the function inputs are already arranged into argument tuples, see *itertools.starmap()*.

- Example usage: Natively multiplying Python lists elementwise,

```
1   from typing import List
2
3   def multiply(x: List, y: List):
4       return x * y
5
6   list_one = [i for i in range(1000)]
7   list_two = [j for j in range(1000, 2000)]
8   result = list(map(multiply, list_one, list_two)) # `map` is a built-
           in function, do not use `(list_one, list_two)` in this case
```

- Example usage: Converting NumPy arrays into PyTorch tensors,

```
1   a = np.array([1, 2, 3, 4])
2   tensor_list = list(map(torch.from_numpy, (a,))) # list containing
           tensor, use of additional brackets necessary
```

- Example usage: Converting NumPy arrays into PyTorch tensors,

```
1   a = np.array([1, 2, 3, 4])
2   b = np.array([5, 6, 7, 8])
3   a, b = map(torch.from_numpy, (a, b)) # tuple unpacking
```

8.4. CPU Brand

Installing py-cpuinfo,

```
1   conda install -c conda-forge py-cpuinfo
```

Obtaining the raw CPU info,

```
1  import cpuinfo
2
3  # Getting detailed CPU information
4  info = cpuinfo.get_cpu_info()
5
6  # If you want to print specific details like the brand, you can do so
   as follows:
7  print(f"CPU Brand: {info['brand_raw']}")
```

8.5. Inheritance

A class inheriting from another class needn't define an `__init__()` function. In that case, the `__init__()` function of the base class will be called.

9. PyTorch

9.1. Leaf Tensors

- If `requires_grad=False`, then the tensor will be leaf by convention. If `requires_grad=True`, then the tensor will be leaf if it was created directly by the user and is **not** the result of an operation, e.g. `.to(device)` when the tensor is on `cpu` and `device="cuda:0"`.
- However, by definition, leaf tensors themselves do not have a **gradient function** `.grad_fn` because they are not the result of a differentiable operation applied to other tensors, i.e. `grad_fn` on such tensors will return `None`. The gradient function in neural network libraries like PyTorch or TensorFlow is associated with tensors that are outputs of differentiable operations.
- The `.grad` attribute on leaf tensors that require gradients, i.e. those for which `requires_grad=True`, stores the gradient computed during backpropagation. (For leaf tensors that have `requires_grad=False`, calling the `.grad` attribute outputs `None`.) Note that for non-leaf tensors, calling `.grad` results in a `UserWarning`, since non-leaf tensors are generally intermediate results in the computation graph, and their gradients are usually not needed once the gradients of the leaf tensors have been obtained. However, there are cases where those gradients are needed, which can be enforced by setting `retain_grad=True` on those tensors,

```
1   x = torch.tensor([1., 2., 3.], requires_grad=True, device=torch.  
    device("cuda:0"))  
2  
3   # fwd pass  
4   y = x**2  
5  
6   # retain gradients  
7   y.retain_grad()  
8  
9   # backward pas  
10  y.sum().backward()  
11  
12  # access gradients  
13  y.grad # `torch.tensor([1., 1., 1.], device="cuda:0")`
```

Note that in the example of this code snippet, doing `y.grad` means that we access the gradient of the scalar loss function `y.sum()` — on which we performed `.backward()`. Correspondingly, doing `x.grad` implies the gradient of the scalar loss function `y.sum()` with respect to `x`.

- In general, it is **not** possible to perform **in-place** operations on leaf tensors for which `requires_grad=True`, since PyTorch dynamically builds a computational graph during

the forward pass, which is used during backpropagation to calculate the gradients. If leaf tensors that have `requires_grad=True` are changed in-place, then the values used during the forward pass are changed, which will affect the gradient calculations in the backward pass. However, note that when no gradients are required for the operations, e.g. when performing parameter updates manually, one can use the context manager with `torch.no_grad()`, in which case in-place operations on leaf tensors **can** be performed, since inside the context manager, `requires_grad=False`.

9.2. Autograd & Backward

- The function `torch.autograd.grad()` computes the gradient. If the gradient of a scalar (loss function) wrt a (weight) matrix is taken, then the output will also be a matrix, where each element corresponds to the partial derivative of the scalar (loss function) wrt to the (weight) matrix element.
- `torch.autograd.grad()` is particularly useful if more direct control over the gradient computation is desired, in particular compared to `.backward()`.
- Note that the default behavior of `.backward()` accumulates gradients in the `.grad` attribute of tensors,

```

1  x = torch.tensor([1., 2., 3.], requires_grad=True, device=torch.
    device("cuda:0"))
2
3  # fwd pass
4  y = 2 * x
5
6  # first backward pass
7  y.sum().backward(retain_graph=True)
8  print(f"Gradients of `x` after first backward pass: {x.grad}") # `
    torch.tensor([2., 2., 2.])`
9
10 # second backward pass
11 y.sum().backward()
12 print(f"Gradients of `x` after second backward pass: {x.grad}") # `
    torch.tensor([4., 4., 4.])`, notice how gradients have
    accumulated

```

However, this behavior can be suppressed by simply zeroing the gradients, i.e. `x.grad.zero_()` — note that `x.grad` returns a tensor, and `<tensor>.zero_()` is a general PyTorch function that sets all elements in-place to 0.

- When using a default iteration loop in PyTorch, `optimizer.zero_grad()` — where `optimizer` is an instance of `torch.optim.Optimizer` — can be put anywhere in the loop *except between* `loss.backward()` and `optimizer.step()`.

9.3. Half-Precision

This tutorial is a great starting point, explaining the advantages of half-precision, i.e. `float16`, training. Implementing this in PYTORCH is possible by following this general recipe, the

documentation of `torch.amp`, in particular the class `torch.autocast`, and some examples can be found [here](#).

A good starting point might also be my [git repo](#), where I implemented a bidirectional LSTM on the MNIST data, which uses the class `torch.autocast`.

9.4. Miscellaneous

Calculating the MSE between two tensors,

```
1 torch.linalg.vector_norm(vt - ut, ord=2,) ** 2 / vt.numel()
```


10. Jax

Try to install via pip first. Only if this doesn't work use conda!

- Putting a Jax array onto a specific device,

```
1  import jax
2  from jax import devices, device_put, numpy as jnp
3
4  x = device_put(jnp.arange(10), device=devices("cpu")[0]) # NOTE: put
   ` [0] `
5  # x = device_put(jnp.arange(10), device=devices("gpu")[0]) # NOTE:
   put ` [0] `
6  print(f"Device: {x.device_buffer.device()}")
```

- Dtype specification,

```
1  x = jnp.array([1, 2, 3], dtype=jnp.float32)
2  print(f"Dtype: {x.dtype}")
```

- Device inference,

```
1  x.device_buffer.device() # x: Jay array
```

- Making a Jax array out of a Python list or a NUMPY array (do not use for tensors),

```
1  from jax import numpy as jnp
2
3  a = jnp.array([1., 2., 3.])
4  b = jnp.array(np.array([1., 2., 3.]))
```

- jit (just-in-time compilation): sets up a function with XLA (extended linear algebra): check out the NB `test__jit-compil.ipynb`. Using jit,

```
1  import jax
2  from jax import numpy as jnp
3
4  @jax.jit
5  def selu(x: jnp.array, lamb: float = 1., alpha: float = 0.):
6      return lamb * jnp.where(x > 0, x, alpha * (jnp.exp(x) - 1.0))
```


A. .bashrc

```
1  ca() {
2      local conda_out="$(conda env list | grep -E "$env_name" | head -n 1 |
3          awk '{print $1}')"
4
5      # check non-emptiness
6      if [ -z "$1" ]; then
7          echo "Usage: ca <env_name>"
8          return 1
9      fi
10
11     # check env existence
12     if [ ! -z "$conda_out" ]; then
13         conda activate "$1"
14     else
15         echo "Conda environment '$env_name' does not exist." # single quotes
16         (') only for display
17         return 1
18     fi
19 }
20
21 # ----- CONDA -----
22
23 # activate conda environment
24 # usage: `ca custom-env-name`
25 ca() {
26     conda activate "$@"
27 }
28
29 # deactivate currently activated conda environment
30 cod() {
31     conda deactivate
32 }
33
34 # List all available conda envs:
35 cel() {
36     conda env list
37 }
38
39 # remove conda environment
40 # usage: `crme ant-migrate-dev`
41 crme() {
```

A. .bashrc

```
41
42 # check number of passed arguments via ` $# `
43 if [[ $# -ne 1 ]]; then
44     echo "NOTE: Exactly one argument needs to be provided"
45 else
46     conda deactivate && conda remove -n "$1" --all -y
47 fi
48
49 }
50
51 # alias for `conda__remove_packages`
52 # usage (e.g.): `crm myenv pkg1 pkg2`
53 crm() {
54     conda__remove_packages "$@"
55 }
56
57 # remove conda packages from environment
58 # usage (e.g.): `conda__remove_packages myenv pkg1 pkg2`
59 conda__remove_packages() {
60
61     # define local variables first
62     local env_name="$1"
63     local conda_out="$(conda env list | grep -E "$env_name" | head -n 1 |
64         awk '{print $1}')"
65
66     # forget first argument (which is saved in `env_name`)
67     shift
68
69     # check non-emptiness
70     if [ -z "$env_name" ]; then
71         echo "Usage: conda__remove_packages <env_name> [package1] [package2]
72             ... [packageN]"
73         return 1
74     fi
75
76     # check env existence
77     if [ ! -z "$conda_out" ]; then
78         conda remove -n "$env_name" "$@" -y
79         echo "Package(s) '$@" removed from environment '$env_name'"
80     else
81         echo "Conda environment '$env_name' does not exist." # single quotes
82         (') only for display
83         return 1
84     fi
85 }
86
87 # ----- AWS -----
```

```

87 # helper function
88 get__profile_endpoint_url() {
89
90     # check if the first argument contains "https://"
91     if [[ "$1" == https://* ]]; then
92         local endpoint_url="$1"
93
94         # if there's a second argument, it's the profile
95         if [ -n "$2" ]; then
96             local profile="$2"
97         fi
98
99     elif [ -n "$1" ]; then
100
101         # if the first argument doesn't contain "https://", it's the profile
102         local profile="$1"
103     fi
104
105     echo "$1 $2"
106
107 }
108
109 # define default vals and update based on provided args
110 update__profile_url() {
111     local endpoint_url="https://kalousis.s3.unige.ch"
112     local profile="default"
113
114     # update `endpoint_url` and `profile` if provided
115     if [[ "$2" == https://* ]]; then
116         read endpoint_url profile <<< $(get__profile_endpoint_url "$2" "$3")
117     else
118         read profile <<< $(get__profile_endpoint_url "$2" "$3") # for `
119             endpoint_url`, default val will be taken
120     fi
121
122     echo "$endpoint_url $profile"
123 }
124
125 # listing
126 # example usages (only bucket name provided):
127 # `lal path`
128 # `lal path default`
129 # `lal path https://kalousis.s3.unige.ch`
130 # `lal path https://kalousis.s3.unige.ch default`
131 lal() {
132     local path="$1"
133
134     read endpoint_url profile <<< $(update__profile_url "$2" "$3")

```

A. .bashrc

```
135     $(which aws) s3 ls s3://" $path" --recursive --endpoint-url "
        $endpoint_url" --profile "$profile"
136 }
137
138 # removing prefixes/files
139 # example usages (only bucket name provided):
140 # `larm path`
141 # `larm path default`
142 # `larm path https://kalousis.s3.unige.ch`
143 # `larm path https://kalousis.s3.unige.ch default`
144 larm() {
145     local path="$1"
146
147     read endpoint_url profile <<< $(update__profile_url "$2" "$3")
148
149     command_output=$(($(which aws) s3 rm s3://" $path" --recursive --
        endpoint-url "$endpoint_url" --profile "$profile")
150
151     if [[ -z "$command_output" ]]; then
152         # no use of `--recursive`, which shouldn't be used for single file
        deletion
153         $(which aws) s3 rm s3://" $path" --endpoint-url "$endpoint_url" --
        profile "$profile"
154     fi
155 }
156
157 # ----- PROTONVPN -----
158
159 p() {
160     protonvpn-cli "$@"
161 }
162
163 # ----- MISCELLANEOUS -----
164
165 # pdflatex
166 pd() {
167     /usr/bin/pdflatex "$@"
168 }
169
170 # convert input notebook to PDF
171 jconv() {
172     jupyter nbconvert --to pdf "$1"
173 }
174
175 # `less` with ANSI escape characters
176 less() {
177     /usr/bin/less -R "$@"
178 }
179
```

```

180 diff() {
181     /usr/bin/diff --color "$@"
182 }
183
184 # shortcut for clearing terminal output
185 c() {
186     clear
187 }
188
189 # shortcuts for exiting terminal
190 q() {
191     exit
192 }
193
194 e() {
195     q
196 }
197
198 # tailscale
199 ts() {
200     tailscale status "$@"
201 }
202
203 # xournalpp (https://github.com/xournalpp/xournalpp)
204 xopp() {
205     xournalpp "$@"
206 }
207
208 # strings comparison
209 # usage (e.g.): `str_diff "blub1" "blub1"` or `str_diff blub1 blub1`
210 # or `str_diff $(echo "hey") $(echo "hey")`
211 # NOTE: exactly two arguments need to be provided
212 str_diff() {
213
214     # check number of passed arguments via ` $# `
215     if [[ $# -ne 2 ]]; then
216         echo "NOTE: Exactly two arguments need to be provided"
217         return 1 # return non-zero exit code to indicate error
218     else
219
220         # compare strings
221         if [[ $1 == $2 ]]; then
222             echo -e "Strings '$1' and '$2' \033[92mmatch\033[0m"
223         else
224             echo -e "Strings '$1' and '$2' do \033[91mNOT\033[0m match"
225         fi
226     fi
227
228 }

```

A. `.bashrc`

```
229
230
231 # ----- DOCKER -----
232 d() {
233     docker "$@"
234 }
235
236 # ----- CHATGPT -----
237
238 # https://github.com/kardolus/chatgpt-cli/tree/main
239 gpt(){
240     chatgpt -i
241 }
242
243 export OPENAI_KEY=[...]
244
245 # ----- ALWAYS EXECUTE -----
246
247 add_bit
```


B. Amazing Programs, Extensions, Plugins & Packages

- <https://etherpad.org/>
- <https://github.com/charmbracelet/glow>
- <https://github.com/0xacx/chatGPT-shell-cli>
- <https://github.com/kardolus/chatgpt-cli/tree/main>

– For setting the right model (cf. here for all available models),

```
1 chatgpt --set-model gpt-4-1106-preview --set-max-tokens 128000
```

– Usage:

```
1 chatgpt -i
```

- <https://tailscale.com/download/>

– Once installation is complete, the command

```
1 sudo tailscale up
```

should be run to login, though this command will also display after installation in the CLI. The signing in should happen via GitHub. To be able to use Tailscale from a new device, it must be added as a device under <https://login.tailscale.com/admin/machines>. Once this is done, open a CLI and type

```
1 ssh name@ip_address # find out <name> and <ip_address> via  
tailscale console  
2 # ssh ellie@100.xx.xxx.xx
```

NOTE that if the file already exists locally, it will be overwritten.

– For file copying (e.g. from the host machine to the currently used machine), do this

```
1 scp name@ip_address:/path/to/remote_file.ext /local/path # find  
out <name> and <ip_address> via tailscale console  
2 # ssh ellie@100.xx.xxx.xx
```

For directory copying,

```
1 scp -r name@ip_address:/path/to/remote_dir /local/path # find  
out <name> and <ip_address> via tailscale console  
2 # ssh ellie@100.xx.xxx.xx
```

B. Amazing Programs, Extensions, Plugins & Packages

- <https://tailscale.com/kb/1080/cli/> (no separate installation necessary, only tailscale needs to be installed)

– Finding out the IPv4 address of the currently active machine,

```
1 tailscale ip -4
```

– Finding out the IPv4 address of another machine connected via the Tailscale network,

```
1 tailscale ip -4 custom-name
2 # tailscale ip -4 ellie
```

- <https://github.com/aws/aws-cli>
- <https://github.com/termcolor/termcolor>
- **LibreOffice dark theme,**

Tools → Options → LibreOffice → Application Colors → Custom Colors → General → Document Background, choose a dark color.

C. Opening Programs from the CLI in Linux

- Opening the settings from CLI,

```
1  gnome-control-center
```

- Opening VSCode from CLI:

```
1  code path_to_file/file_name.ext
```

If a VSCode editor is already open, use the *-n* flag to open the file in a new editor:

```
1  code -n path_to_file/file_name.ext
```

A folder can also be opened directly:

```
1  code path_to_dir
```

Listing C.1: Opening VSCode dir from CLI

- Opening LibreOffice from CLI:

```
1  libreoffice --writer path_to_dir/filename.odt
```

- Opening an image via the CLI:

```
1  eog /path/to/your/image.jpg
```


D. VSCode

D.1. Recommended Extensions

- <https://marketplace.visualstudio.com/items?itemName=ms-vscode-remote.vscode-remote>
- <https://marketplace.visualstudio.com/items?itemName=Gruntfuggly.todo-tree>

D.2. Debugging

- Stepping into external code with Python debugger, tutorial [here](#)
- Creating a JSON file, [here](#) some instructions

D.3. settings.json

Opening the file,

1. press `Ctrl + Shift + P` to the Command Palette,
2. type `Open User Settings (JSON)` and select it to open the `settings.json` file.

D.4. Fix Unresolved Python Imports

- If you run a docker container where a conda environment is installed (with packages that you do not have locally), then VSCode will show those imports as unresolved. To fix this, open the `settings.json` file, cf. App. D.3, and add the following setting:

```
1  "python.analysis.diagnosticSeverityOverrides": {  
2    "reportMissingImports": "none"  
3  }
```

Incorporating this into the `settings.json` file is shown in App. D.6.

- Note that if you have an SSH connection to another machine going on, e.g. in the Remote Development extension, putting the above lines into the `settings.json` file will not have an immediate effect, for this the SSH connection needs to be restarted.

D.5. Opening a Duplicate Workspace

1. press `Ctrl + Shift + P` to open the Command Palette,
2. then type `Workspaces: Duplicate As Workspace in New Window`

D.6. settings.json

Contents of settings.json,

```
1  {
2    "workbench.colorTheme": "Default Dark Modern",
3    "telemetry.telemetryLevel": "off",
4    "editor.wordWrap": "wordWrapColumn",
5    "editor.wordWrapColumn": 79,
6    "workbench.editor.enablePreview": false,
7    "gitlens.telemetry.enabled": false,
8    "notebook.lineNumbers": "on",
9    "explorer.confirmDragAndDrop": false,
10   "window.zoomLevel": 1,
11   "python.analysis.diagnosticSeverityOverrides": {
12     "reportMissingImports": "none"
13   },
14   "todo-tree.general.tags": [
15     "BUG",
16     "HACK",
17     "FIXME",
18     "TODO",
19     "NOTE",
20     "XXX",
21     "[ ]",
22     "[x]"
23   ],
24
25   "files.associations": {"*.log": "plaintext"},
26   "[plaintext]": {"editor.wordWrap": "off"},
27   "[shellscript]": {"editor.wordWrap": "off"},
28   "workbench.editor.tabSizing": "shrink"
29 }
```